D04 · Category Strategy & Spend Analytics

01 · Spend analytics & categorization

Supply Chain Reference Architecture v2.2 · mrsupplychain.ai

1 WHAT IT IS

Spend analytics is the data foundation under any category strategy — extracting AP and PO data from the ERP, normalizing supplier names, mapping spend to a hierarchical taxonomy (UNSPSC or custom), and producing the multidimensional spend cube that procurement decisions run on. It is the unglamorous first step that determines whether everything downstream is grounded in reality or in storytelling.

2 WHY IT MATTERS

Spend analytics is the discipline that procurement teams most often skip — and most consistently regret. **The typical large enterprise has between 25% and 45% of its supplier base duplicated across name variants** ("ACME Corp", "Acme Corporation", "ACME CORP."), and between 30% and 60% of its spend mapped to the wrong category or to no category at all **(Hackett Group, 2024 — procurement benchmarking)**. The asymmetry is severe: a category strategy built on a 30%-wrong spend cube produces decisions that are 30% wrong, and the downstream cost (wrong consolidation targets, wrong negotiation positions, wrong supplier rationalization) compounds for years. **The economic case is documented across multiple sources** — clean spend data is the precondition for procurement value capture, not an output of it **(Kearney, 2023 — Assessment of Excellence in Procurement study)**.

3 HOW IT IS DONE

1. **Extract from the source systems and accept that the data is dirty.** Pull AP transactions, PO data, and contract data from the ERP — knowing that 30–60% will need cleansing. The first extraction is never the analysis-ready dataset; pretending otherwise hides the data quality problem instead of solving it.
2. **Normalize supplier names mechanically — not manually.** Run fuzzy matching (Levenshtein, soundex), parent-subsidiary resolution via DUNS or tax IDs, and supplier consolidation rules. Manual cleansing scales linearly with supplier count; mechanical cleansing scales logarithmically. A 12,000-supplier base needs automation.
3. **Apply a taxonomy — UNSPSC, NAICS, or custom — consistently across all transactions.** The taxonomy choice is less important than the consistency. UNSPSC is the international standard with 8-digit category resolution; many large operations use a custom hybrid that maps cleanly to their internal category management structure.
4. **Build the spend cube along the three core dimensions — supplier, category, business unit — plus geography and time.** The cube is the analytical artifact that procurement leadership operates on. It must support drill-down from total spend to single transaction, and roll-up from single transaction to corporate view.
5. **Surface the off-contract spend explicitly.** Match every transaction against the active contract base. The percentage of spend running outside negotiated contracts (maverick buying) is the single most actionable number to come out of the spend cube — and the one most plants do not measure.
6. **Refresh the cube monthly, not annually.** Spend patterns shift as the business evolves; a category that was leverage spend last year may now be strategic. The monthly refresh keeps category management responsive; the annual refresh ensures decisions are always at least 11 months stale.

WORKED EXAMPLE — ILLUSTRATIVE

A diversified manufacturer with $1.2B annual spend across 4 business units pulls 18 months of AP data — **340,000 transactions across 8,400 distinct supplier names**. Fuzzy matching and DUNS resolution collapses the supplier list to **6,200 unique entities** (26% reduction from name variants). UNSPSC categorization reveals that **$340M (28%) of spend was previously uncategorized or in "miscellaneous"**, hiding the actual category mix. The off-contract spend analysis surfaces **$42M of maverick buying** — purchases made outside the 247 active master agreements, mostly in indirect categories (IT, MRO, facilities). The cube becomes the single source of truth for the upcoming category review; the data cleansing project itself pays back in 5 months from off-contract spend recapture alone.

IN PRACTICE

Normalize supplier names mechanically

Build a deterministic dedup pipeline — fuzzy match plus a manual review queue — and run it before every cube refresh. Manual cleansing is unrepeatable and decays the moment a new supplier is onboarded; mechanical normalization is the only approach that survives the next monthly cycle.

Pick one taxonomy and apply it consistently

Choose UNSPSC, eClass, or a custom hierarchy and commit. The worst outcome is two taxonomies live in different reports. The taxonomy choice matters less than the discipline of applying it consistently — analysts can adapt to any tree, they cannot adapt to a moving one.

Refresh the cube monthly, not on demand

A spend cube rebuilt only when someone asks is a cube that is always stale at the moment of the question. Monthly cadence is the minimum that keeps the data fresh enough for off-contract spend, consolidation opportunities, and category drift to surface in time to act on them.

Surface off-contract spend explicitly

Build a single dashboard line per strategic category that quantifies spend bypassing the awarded supplier or contract. The number is almost always larger than expected; surfacing it explicitly is the precondition for closing the leakage — usually worth more than the next round of negotiated savings.

WHAT YOU WOULD USE

SPEND ANALYTICS PLATFORMS

Coupa Spend Analysis / SAP Ariba Spend Visibility — Enterprise spend analytics modules integrated into the broader procurement suite — strongest when the operation is already standardized on one of these stacks **(Coupa — vendor; SAP — vendor)**.

Sievo, Suplari, GEP Smart — Specialized spend analytics platforms with strong supplier normalization engines and pre-built dashboards — faster to deploy than enterprise suites, often the right entry point **(Sievo — vendor; GEP — vendor)**.

Consultancy-driven analysis (Kearney, McKinsey, BCG) — When the spend analytics capability does not exist internally and the operation needs the cube built once with knowledge transfer — strong for one-time strategic resets, less suited for ongoing operations.

Custom Python + Tableau / Power BI stack — For operations with the in-house data engineering capability — pandas for cleansing, fuzzy matching libraries (FuzzyWuzzy, RapidFuzz) for normalization, Tableau or Power BI for the visualization layer. Highest ceiling, highest commitment.

4 SO WHAT

The category strategy that is built on dirty spend data is fiction with a budget. Extract from source systems and accept the data is dirty, normalize supplier names mechanically, apply a consistent taxonomy, build the spend cube along the right dimensions, surface off-contract spend explicitly, and refresh monthly. The plants that do this well buy back the precondition for every other procurement decision; the plants that skip it spend the next three years debating which version of the numbers is the right one.

From the SCRA v2.2 reference architecture · mrsupplychain.ai · © 2026 Ismael López · Mr. Supply Chain® México